

16.10.2023

Digital Image Processing

Image Compression

Makarand Tapaswi



Center for Visual Information Technology (CVIT), IIT Hyderabad



Most slides borrowed from Ravi Kiran @CVIT!

14.10.2022

Digital Image Processing (CSE/ECE 478)

Lecture-17: Image Compression

Ravi Kiran



Center for Visual Information Technology (CVIT), IIT Hyderabad



Some slides borrowed from Vineet Gandhi @CVIT!

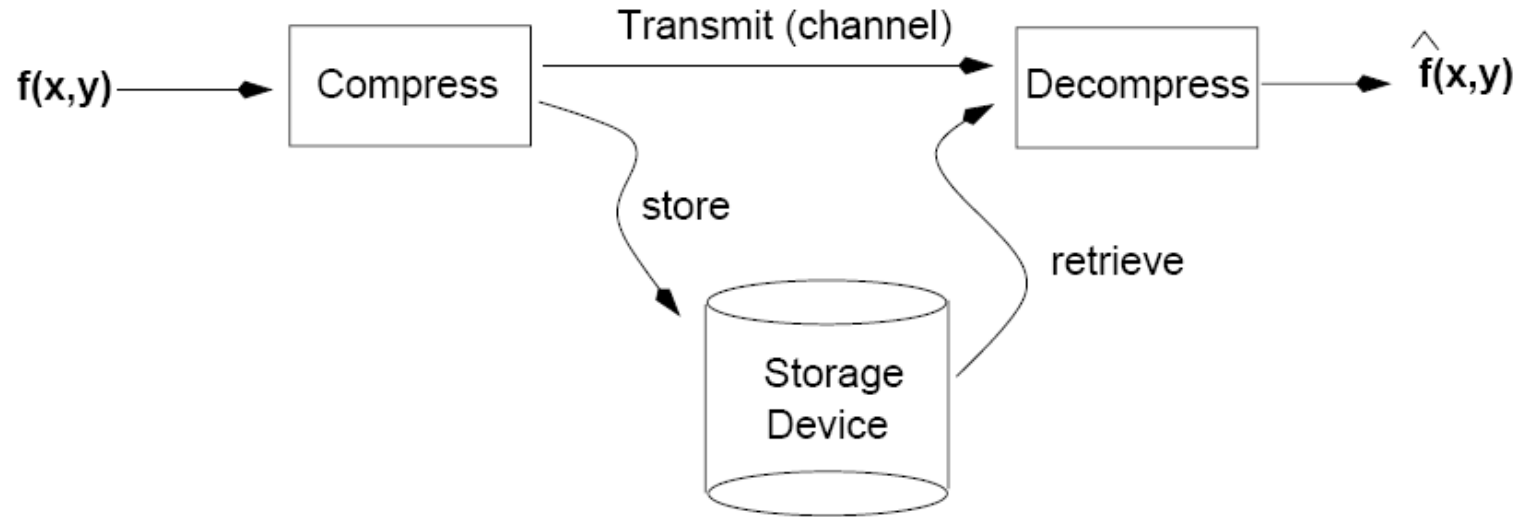


Data Compression

- Data compression aims to reduce the amount of data while preserving as much **information** as possible.

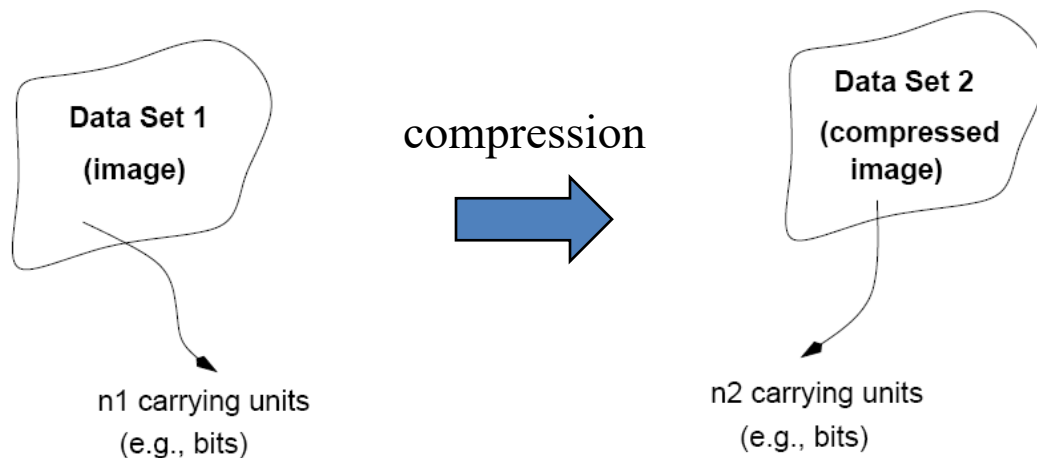
Image Compression

- Goal: Reduce amount of data required to represent a digital image (c.f. signal).



.. By exploiting redundancies in image data

Compression Ratio



Compression ratio: $C_R = \frac{n_1}{n_2}$

Relevant Data Redundancy

$$R_D = 1 - \frac{1}{C_R}$$

Example:

$$\text{If } C_R = \frac{10}{1}, \text{ then } R_D = 1 - \frac{1}{10} = 0.9$$

(90% of the data in dataset 1 is redundant)

$$\text{if } n_2 = n_1, \text{ then } C_R=1, R_D=0$$

$$\text{if } n_2 \ll n_1, \text{ then } C_R \rightarrow \infty, R_D \rightarrow 1$$

Motivation

- Consider a 2 hour, full HD video (1920x1080 resolution@30fps)
- The storage space required for raw data is $1920 \times 1080 \times 24 \text{ bits} = 6.22 \text{ MB}$
- Space required per frame is $6.22 \times 10^6 \text{ bits} = 186.6 \text{ MB}$
- Space required for entire movie is $186.6 \times 30 \times 2 \times 60 \times 60 \text{ bits} = 1.34 \times 10^{12} \text{ bytes} = 1.34 \text{ TB}$
- To put it on a 25 GB blu ray disc: required compression factor = **53.7**
- To put it in a 1GB mp4 file: required compression factor = **1343**



Diversion: bits and bytes, kilo vs kibi

- Previous convention: divide by 1024
- New convention: divide by 1000
- SI unit: kB (kilo = 1000)
- This is what causes a 4TB disk to have 3.7TiB!
- Also, note: B is not the same as b

Types of Redundancy

(1) Coding Redundancy

(2) Spatial/Temporal Redundancy

(3) Psychovisual Redundancy

- Image compression attempts to reduce one or more of these redundancy types.

Types of Redundancy

(1) Coding Redundancy

(2) Spatial/Temporal Redundancy

(3) Psychovisual Redundancy

- Image compression attempts to reduce one or more of these redundancy types.

Coding - Definitions

- **Code:** a list of symbols (letters, numbers, bits etc.)
- **Code word:** a sequence of symbols used to represent some information (e.g., gray levels).
- **Code word length:** number of symbols in a code word.

Example: (binary code, symbols: 0,1, length: 3)

0: 000	4: 100
1: 001	5: 101
2: 010	6: 110
3: 011	7: 111

Coding - Definitions (cont'd)

N x M image

r_k: k-th gray level

l(r_k): # of bits for r_k

P(r_k): probability of r_k



Discrete r.v. in the range [0,L-1]

Average # of bits: $L_{avg} = E(l(r_k)) = \sum_{k=0}^{L-1} l(r_k)P(r_k)$

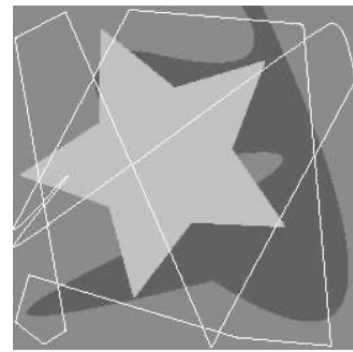
Total # of bits: NML_{avg}



a b c

FIGURE 8.1 Computer generated $256 \times 256 \times 8$ bit images with (a) coding redundancy, (b) spatial redundancy, and (c) irrelevant information. (Each was designed to demonstrate one principal redundancy but may exhibit others as well.)

Coding Redundancy



- Case 1: $l(r_k) = \text{constant length}$

Example:

r_k	$p_r(r_k)$	Code 1	$l_1(r_k)$
$r_0 = 0$	0.19	000	3
$r_1 = 1/7$	0.25	001	3
$r_2 = 2/7$	0.21	010	3
$r_3 = 3/7$	0.16	011	3
$r_4 = 4/7$	0.08	100	3
$r_5 = 5/7$	0.06	101	3
$r_6 = 6/7$	0.03	110	3
$r_7 = 1$	0.02	111	3

$$\text{Average \# of bits: } L_{avg} = E(l(r_k)) = \sum_{k=0}^{L-1} l(r_k)P(r_k)$$

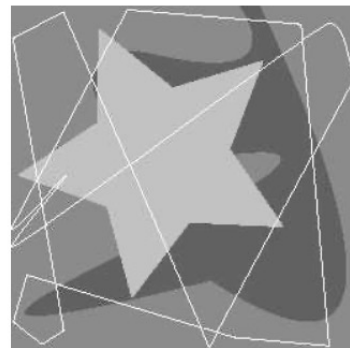
$$\text{Total \# of bits: } NML_{avg}$$

Assume an image with $L = 8$

$$\text{Assume } l(r_k) = 3, \quad L_{avg} = \sum_{k=0}^7 3P(r_k) = 3 \sum_{k=0}^7 P(r_k) = 3 \text{ bits}$$

Total number of bits: $3NM$

Coding Redundancy (cont'd)



- Case 2: $l(r_k) = \text{variable length}$

Table 6.1 Variable-Length Coding Example variable length

r_k	$p_r(r_k)$	Code 1	$l_1(r_k)$	Code 2	$l_2(r_k)$
$r_0 = 0$	0.19	000	3	11	2
$r_1 = 1/7$	0.25	001	3	01	2
$r_2 = 2/7$	0.21	010	3	10	2
$r_3 = 3/7$	0.16	011	3	001	3
$r_4 = 4/7$	0.08	100	3	0001	4
$r_5 = 5/7$	0.06	101	3	00001	5
$r_6 = 6/7$	0.03	110	3	000001	6
$r_7 = 1$	0.02	111	3	000000	6

$$C_R = \frac{n_1}{n_2}$$

$$L_{avg} = \sum_{k=0}^7 l(r_k)P(r_k) = 2.7 \text{ bits}$$

Total number of bits: $2.7NM$

$$C_R = \frac{3}{2.7} = 1.11 \text{ (about 10\%)}$$

$$R_D = 1 - \frac{1}{1.11} = 0.099$$

Types of Redundancy

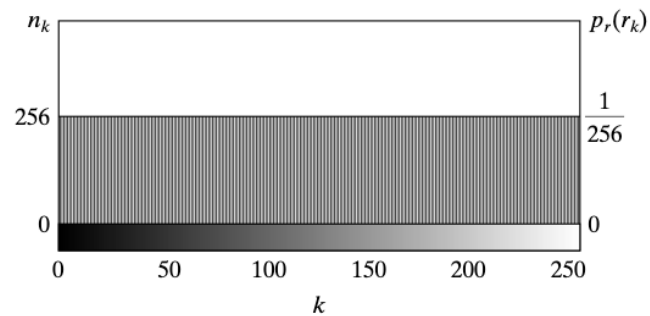
(1) Coding Redundancy

(2) Spatial/Temporal Redundancy

(3) Psychovisual Redundancy

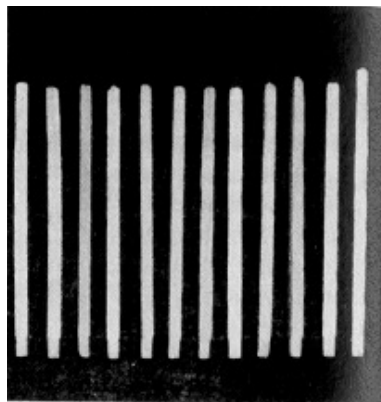
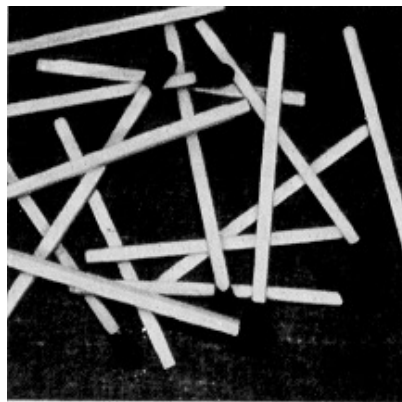
- Image compression attempts to reduce one or more of these redundancy types.

Spatial Redundancy



Spatial redundancy

- Interpixel redundancy exists → pixel values are correlated
- i.e., a pixel value can be reasonably predicted by its neighbors



histograms

auto-correlation

$$f(x) \circ g(x) = \int_{-\infty}^{\infty} f(x)g(x+a)da$$

auto-correlation: $f(x)=g(x)$




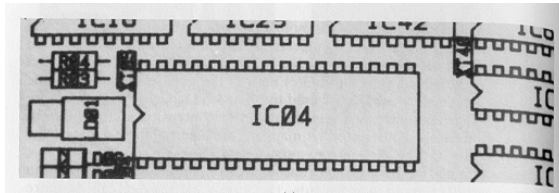
Figure 6.2 Two images and their gray-level histograms and normalized autocorrelation

Interpixel redundancy (cont'd)

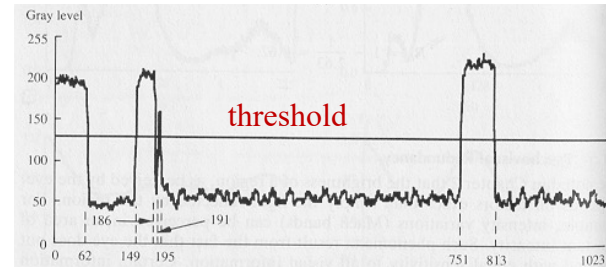
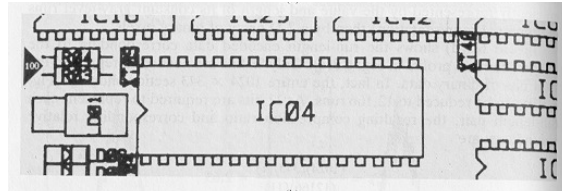
- To reduce interpixel redundancy, some kind of transformation must be applied on the data (e.g., thresholding, DFT, DWT)

Example:

original



thresholded



110000.....11.....000.....

Run-length encoding:

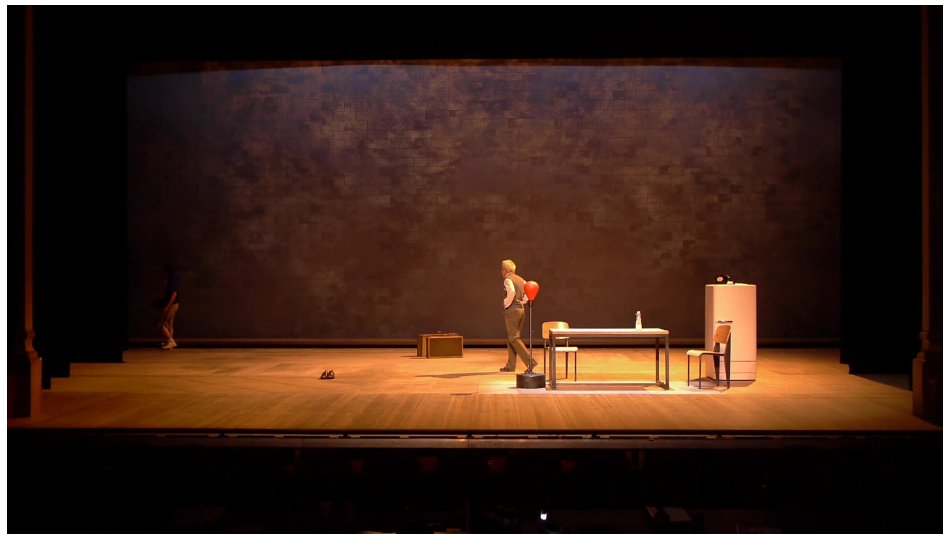
(1,63) (0,87) (1,37) (0,5) (1,4) (0, 556) (1,62) (0,210)

Using 11 bits/pair:

(1+10) bits/pair

88 bits are required (compared to 1024 !!)

Spatial and temporal redundancy



frame t



frame t+1

Spatial and temporal redundancy



Types of Redundancy

(1) Coding Redundancy

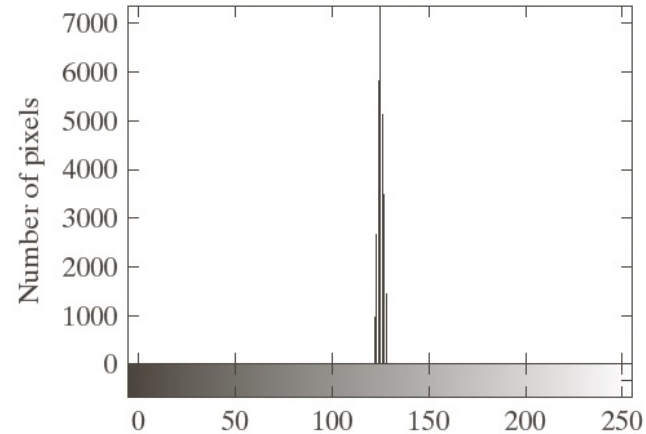
(2) Spatial/Temporal Redundancy

(3) Psychovisual Redundancy

- Image compression attempts to reduce one or more of these redundancy types.

Irrelevant information or perceptual redundancy

- Not all visual information is perceived by eye/brain, so throw away those that are not



Psychovisual redundancy (cont'd)

Example: quantization

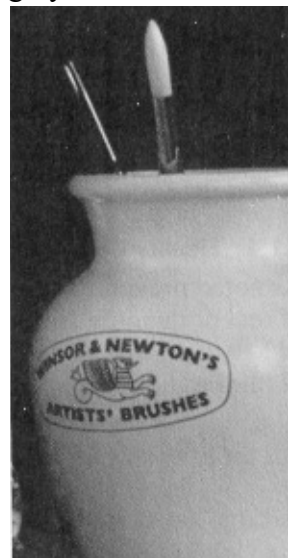
256 gray levels



16 gray levels



16 gray levels + random noise



$$C=8/4 = 2:1$$

add a small pseudo-random number
to each pixel prior to quantization

Information theory

- Basic Premise: Generation of information can be treated as a probabilistic process defined over symbols.
- Symbol - carrier of information
- Consider a symbol with an occurrence probability p .
- The amount of information contained in the symbol is defined as:

$$I = \log_2 \frac{1}{p} \text{ bits} \quad \text{or} \quad I = -\log_2 p$$

Information theory: Entropy

- Consider a source that contains L possible symbols $\{s, i=0,1,2,\dots,L-1\}$
- With corresponding occurrence probabilities defined as $\{p_i, i=0,1,2,\dots,L-1\}$

- **Entropy**

$$H = - \sum_{i=0}^{L-1} p_i \log_2 p_i$$

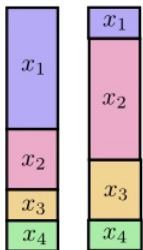
r_k	$p_r(r_k)$	Code 1	$l_1(r_k)$	Code 2	$l_2(r_k)$
$r_{87} = 87$	0.25	01010111	8	01	2
$r_{128} = 128$	0.47	10000000	8	1	1
$r_{186} = 186$	0.25	11000100	8	000	3
$r_{255} = 255$	0.03	11111111	8	001	3
r_k for $k \neq 87, 128, 186, 255$	0	—	8	—	0

$$\log(0.47) = -1.09$$

$$\log(0.03) = -5.06$$

Slight Detour: Cross-entropy

$p(x)$ $q(x)$



Cross-Entropy: $H_p(q)$

Average Length
of message from $q(x)$
using code for $p(x)$.

$$H = - \sum_{i=0}^{L-1} p_i \log_2 p_i$$

The diagram illustrates the cross-entropy formula $D(\hat{\mathbf{y}}, \mathbf{y}) = - \sum_j y_j \ln \hat{y}_j$. It shows two vectors: $\hat{\mathbf{y}} = \begin{bmatrix} 0.1 \\ 0.5 \\ 0.4 \end{bmatrix}$ and $\mathbf{y} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$. A red arrow points from the top element of $\hat{\mathbf{y}}$ to the top element of \mathbf{y} , and a blue arrow points from the middle element of \mathbf{y} to the middle element of $\hat{\mathbf{y}}$. The formula is written in the center, with the terms y_j and \hat{y}_j highlighted in blue and red respectively.

Information theory: Shannon's theorem

- Shannon's lossless source coding theorem: For a discrete, memoryless, stationary information source, the minimum bit rate required to encode a symbol on average is equal to the entropy of the source.
- In other words: we can't do better than the entropy
- Let's understand with an example

r_k	$p_r(r_k)$	Code 1	$l_1(r_k)$	Code 2	$l_2(r_k)$
$r_{87} = 87$	0.25	01010111	8	01	2
$r_{128} = 128$	0.47	10000000	8	1	1
$r_{186} = 186$	0.25	11000100	8	000	3
$r_{255} = 255$	0.03	11111111	8	001	3
r_k for $k \neq 87, 128, 186, 255$	0	—	8	—	0

Types of Redundancy

(1) Coding Redundancy

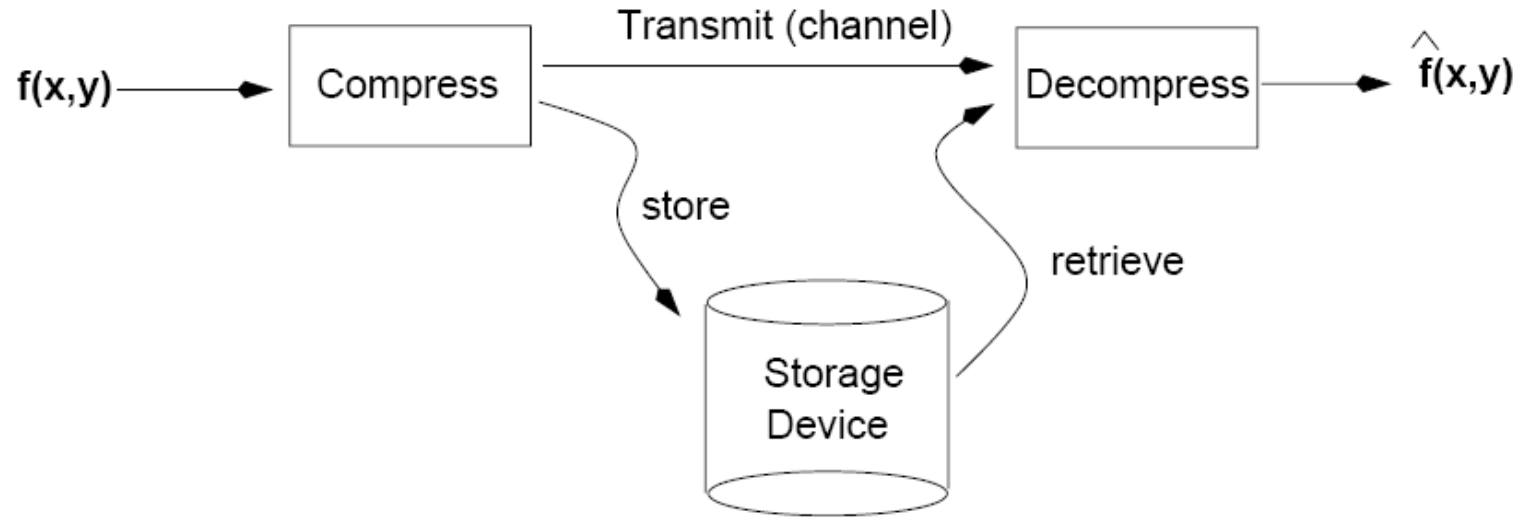
(2) Spatial/Temporal Redundancy

(3) Psychovisual Redundancy

- Image compression attempts to reduce one or more of these redundancy types.

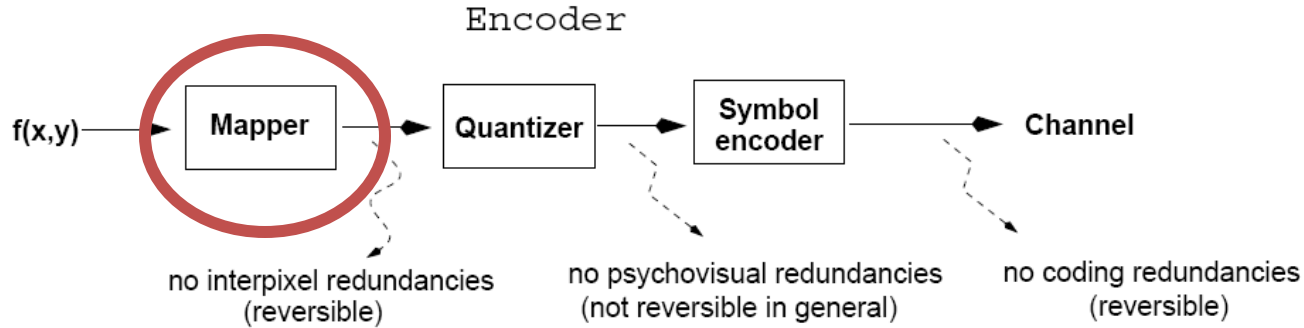
Image Compression

- Goal: Reduce amount of data required to represent a digital image (c.f. signal).



.. By exploiting redundancies in image data

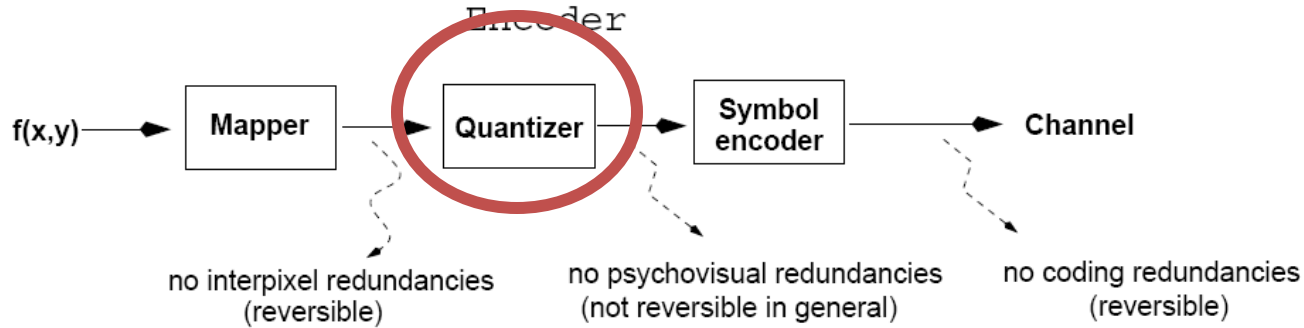
Image Compression Model



-

- **Mapper:** transforms data to account for interpixel redundancies.

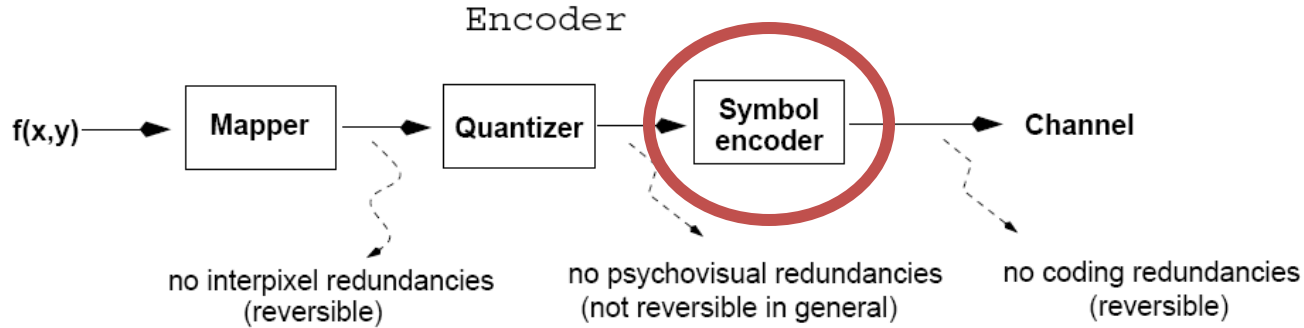
Image Compression Model (cont'd)



-

- **Quantizer:** quantizes the data to account for psychovisual redundancies.


Image Compression Model (cont'd)



-


- **Symbol encoder:** encodes the data to account for coding redundancies.

Image Compression Models (cont'd)



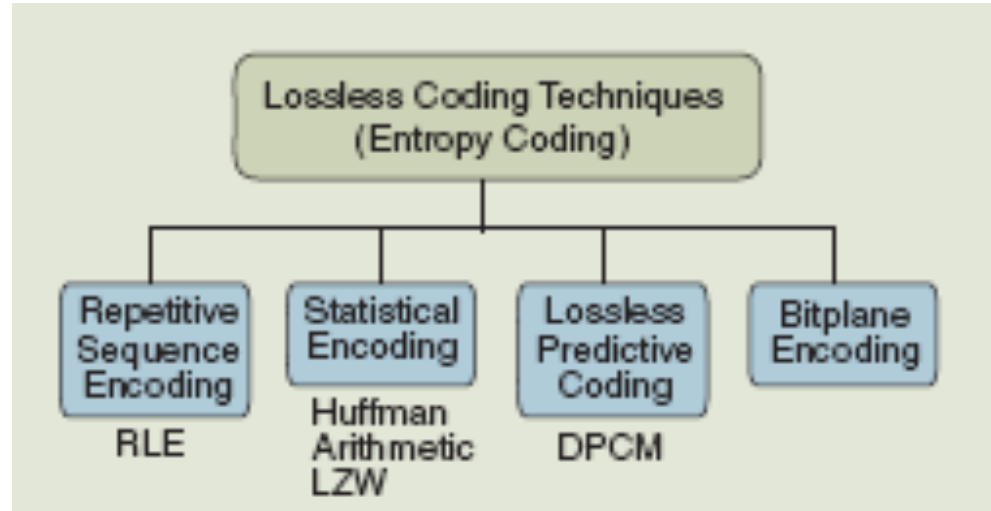
- The decoder applies the inverse steps.
- Note that quantization is **irreversible** in general.

Lossless Compression



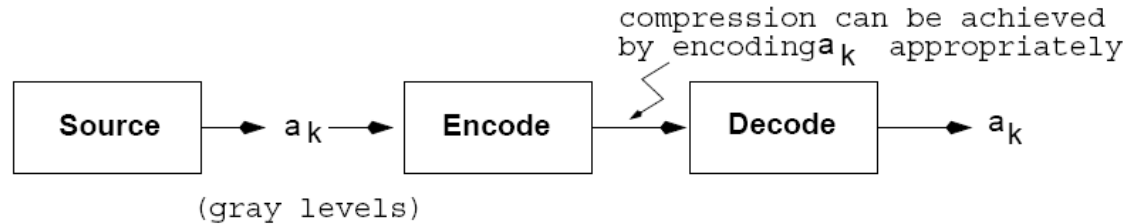
$$e(x, y) = \hat{f}(x, y) - f(x, y) = 0$$

Taxonomy of Lossless Methods



Huffman Coding

(addresses coding redundancy)



- A **variable-length coding** technique.
- Source symbols are encoded **one** at a time!
 - There is a **one-to-one correspondence** between source symbols and code words.
- **Optimal code** - minimizes code word length per source symbol.

Huffman Coding (cont'd)

- Forward Pass

1. Sort probabilities per symbol
2. Combine the lowest two probabilities
3. Repeat *Step2* until only two probabilities remain.

Original source		Source reduction			
Symbol	Probability	1	2	3	4
a_2	0.4	0.4	0.4	0.4	0.6
a_6	0.3	0.3	0.3	0.3	
a_1	0.1	0.1	0.2	0.3	0.4
a_4	0.1	0.1			
a_3	0.06	0.1	0.1	0.3	0.4
a_5	0.04		0.1		

- Backward Pass

Assign code symbols going backwards

Original source			Source reduction			
Sym.	Prob.	Code	1	2	3	4
a_2	0.4	1	0.4 1	0.4 1	0.4 1	0.6 0
a_6	0.3	00	0.3 00	0.3 00	0.3 00	
a_1	0.1	011	0.1 011	0.2 010	0.3 01	0.4 1
a_4	0.1	0100	0.1 0100			
a_3	0.06	01010	0.1 0101	0.1 011	0.1 010	0.1 011
a_5	0.04	01011				

Huffman Coding (cont'd)

- L_{avg} assuming Huffman coding:

$$L_{avg} = E(l(a_k)) = \sum_{k=1}^6 l(a_k)P(a_k) =$$

$$3 \times 0.1 + 1 \times 0.4 + 5 \times 0.06 + 4 \times 0.1 + 5 \times 0.04 + 2 \times 0.3 = 2.2 \text{ bits/symbol}$$

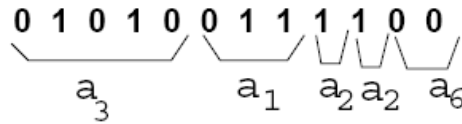
6 symbols, we need a 3-bit code

- $(a_1: 000, a_2: 001, a_3: 010, a_4: 011, a_5: 100, a_6: 101)$

$$L_{avg} = \sum_{k=1}^6 l(a_k)P(a_k) = \sum_{k=1}^6 3P(a_k) = 3 \sum_{k=1}^6 P(a_k) = 3 \text{ bits/symbol}$$

Huffman Coding/Decoding

- Coding/Decoding can be implemented using a **look-up table**.
- Decoding can be done unambiguously.



Original source			Source reduction			
Sym.	Prob.	Code	1	2	3	4
a_2	0.4	1	0.4 1	0.4 1	0.4 1	0.6 0
a_6	0.3	00	0.3 00	0.3 00	0.3 00	0.4 1
a_1	0.1	011	0.1 011	0.2 010	0.3 01	
a_4	0.1	0100	0.1 0100	0.1 011		
a_3	0.06	01010	0.1 0101			
a_5	0.04	01011				

Original source		
Sym.	Prob.	Code
a_2	0.4	1
a_6	0.3	00
a_1	0.1	011
a_4	0.1	0100
a_3	0.06	01010
a_5	0.04	01011

Arithmetic (or Range) Coding (addresses coding redundancy)

- The main weakness of Huffman coding is that it encodes source symbols **one** at a time.
- Arithmetic coding encodes **sequences** of source symbols together.
 - There is **no** one-to-one correspondence between source symbols and code words.
- Slower than Huffman coding but can achieve better compression.

Arithmetic Coding (cont'd)

- A sequence of source symbols is assigned to a sub-interval in $[0,1)$ which can be represented by an arithmetic code, e.g.:

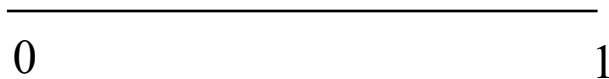


- Start with the interval $[0, 1)$; a sub-interval is chosen to represent the message which becomes smaller and smaller as the number of symbols in the message increases.


Arithmetic Coding (cont'd)

Encode message: $\alpha_1 \alpha_2 \alpha_3 \alpha_3 \alpha_4$

1) Start with interval $[0, 1)$



2) Subdivide $[0, 1)$ based on the probabilities of α_i




Source Symbol	Probability
a_1	0.2
a_2	0.2
a_3	0.4
a_4	0.2


Initial Subinterval
$[0.0, 0.2)$
$[0.2, 0.4)$
$[0.4, 0.8)$
$[0.8, 1.0)$

3) Update interval by processing source symbols

Example

Source Symbol	Probability	Initial Subinterval
a_1	0.2	$[0.0, 0.2)$
a_2	0.2	$[0.2, 0.4)$
a_3	0.4	$[0.4, 0.8)$
a_4	0.2	$[0.8, 1.0)$



Encode
 $\alpha_1 \alpha_2 \alpha_3 \alpha_3 \alpha_4$

 $[0.06752, 0.0688)$
 or
 0.068
 (must be inside sub-interval)


Example (cont'd)

- The message $\alpha_1 \alpha_2 \alpha_3 \alpha_3 \alpha_4$ is encoded using 3 decimal digits or $3/5 = 0.6$ decimal digits per source symbol.
- The entropy of this message is: $H = - \sum_{k=0}^3 P(r_k) \log(P(r_k))$

$$-(3 \times 0.2 \log_{10}(0.2) + 0.4 \log_{10}(0.4)) = 0.5786 \text{ digits/symbol}$$

Note: finite precision arithmetic might cause problems due to truncations!

Arithmetic Decoding



Decode 0.572
(code length=4)



$\alpha_3 \alpha_3 \alpha_1 \alpha_2 \alpha_4$

Reference

- Ch 8, G&W textbook